

Support Vector Machines

Based on ESL and papers by Vladimir Vapnik, Trevor Hastie, Saharon Rosset, Rob Tibshirani, Ji Zhu


Outline

- Optimal separating hyperplanes and relaxations
- SVMs and kernel inner-products
- SVM as a function estimation problem
- LARS- style algorithm for SVMs

Maximum Margin Classifier


Vapnik(1995)

$x_i \in \mathbb{R}^p$, $y_i \in \{-1, 1\}$



$$\begin{aligned} & \max_{\beta, \beta_0, \|\beta\|=1} C \\ \text{subject to} \quad & y_i(x_i^T \beta + \beta_0) \geq C, \quad i = 1, \dots, N. \end{aligned}$$

Overlapping Classes



$$\begin{aligned} & \max_{\beta, \beta_0, \|\beta\|=1} C \\ & \text{subject to } y_i(x_i^T \beta + \beta_0) \geq C(1 - \xi_i), \quad \xi_i \geq 0, \quad \sum_i \xi_i \leq B \end{aligned}$$

Equivalent form of problem


Define $C = 1/||\beta||$ and drop norm constraint on β , as in ESL sec 4.2:

$$\min ||\beta||$$

subject to $y_i(x_i^T \beta + \beta_0) \geq 1 - \xi_i$, $\xi_i \geq 0$, $\sum_i \xi_i \leq B$

This is the original form given by Vapnik; we find it confusing due to the fixed scale “1” in the constraint.

Example



Fitted function is $\hat{f}(x) = x^T \hat{\beta} + \hat{\beta}_0$

Resulting classifier is $\hat{G}(x) = \text{sign}[\hat{f}(x)]$

Quadratic Programming Solution

After a lot of *stuff* we arrive at a Lagrange dual

$$L_D = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{i'=1}^N \alpha_i \alpha_{i'} y_i y_{i'} x_i^T x_{i'}$$

which we maximize subject to constraints (involving B as well).

The solution is expressed in terms of fitted Lagrange multipliers $\hat{\alpha}_i$:

$$\hat{\beta} = \sum_{i=1}^N \hat{\alpha}_i y_i x_i$$

Some fraction of $\hat{\alpha}_i$ are exactly zero (from KKT conditions); the x_i for which $\hat{\alpha}_i > 0$ are called **support points** \mathcal{S} .

$$\hat{f}(x) = x^T \hat{\beta} + \hat{\beta}^0 = \sum_{i \in \mathcal{S}} \hat{\alpha}_i y_i x^T x_i + \hat{\beta}^0$$


Microarray example

16,063 genes; 144 training, 54 test samples, 14 classes

Methods	CV errors (SE) Out of 144	Test errors Out of 54	Number of Genes Used
<hr/>			
1. Nearest shrunken centroids	35 (5.0)	17	6,520
2. L_2 -penalized discriminant analysis	25 (4.1)	12	16,063
3. Support vector classifier	26 (4.2)	14	16,063
4. Lasso regression (one vs all)	30.7 (1.8)	12.5	1,429
5. k -nearest neighbors	41 (4.6)	26	16,063
6. L_2 -penalized multinomial	26 (4.2)	15	16,063
7. L_1 -penalized multinomial	17 (2.8)	13	269
8. Elastic-net penalized multinomial	22 (3.7)	11.8	384

Flexible Classifiers

SVM - Degree-4 Polynomial in Feature Space



Enlarge the feature space via basis expansions, e.g. polynomials of total degree 4. $h(x) = (h_1(x), h_2(x), \dots, h_M(x))$

$$\hat{f}(x) = h(x)^T \hat{\beta} + \hat{\beta}_0$$

The kernel trick

- Consider the ridge regression prediction

$$\hat{\mathbf{y}} = \mathbf{X}(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}_p)^{-1} \mathbf{X}^T \mathbf{y}$$

If p is large, computations with $\mathbf{X}^T \mathbf{X} + \lambda I$ ($p \times p$ matrix) may be daunting. Instead write this as

$$\hat{\mathbf{y}} = (\mathbf{X}\mathbf{X}^T + \lambda \mathbf{I}_N)^{-1} \mathbf{X}\mathbf{X}^T \mathbf{y}$$

- matrix is now only $N \times N$, and if N is small, this is much simpler computationally.
- note that we need only to compute inner products of the observations, to compute $\mathbf{X}\mathbf{X}^T$. This argument also applies if X represents not the original features but transformations of them. Hence we only need to define inner products in the transformed space; we don't need to actually transform the features!

SVM and Kernels

$$L_D = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{i'=1}^N \alpha_i \alpha_{i'} y_i y_{i'} \langle h(x_i), h(x_{i'}) \rangle$$

$$\begin{aligned} f(x) &= h(x)^T \beta + \beta_0 \\ &= \sum_{i=1}^N \alpha_i y_i \langle h(x), h(x_i) \rangle + \beta_0. \end{aligned}$$

L_D and solution $f(x)$ involve $h(x)$ only through inner-products

$$K(x, x') = \langle h(x), h(x') \rangle$$

Given a suitable positive kernel $K(x, x')$, don't need $h(x)$ at all!

$$\hat{f}(x) = \sum_{i \in \mathcal{S}} \hat{\alpha}_i y_i K(x, x_i) + \hat{\beta}_0$$

Popular Kernels

$K(x, x')$ is a symmetric, positive (semi-)definite function.

*d*th deg. poly.: $K(x, x') = (1 + \langle x, x' \rangle)^d$

radial basis: $K(x, x') = \exp(-\|x - x'\|^2 / c)$

Example: 2nd degree polynomial in \mathbb{R}^2 .

$$\begin{aligned} K(x, x') &= (1 + \langle x, x' \rangle)^2 \\ &= (1 + x_1 x'_1 + x_2 x'_2)^2 \\ &= 1 + 2x_1 x'_1 + 2x_2 x'_2 + (x_1 x'_1)^2 + (x_2 x'_2)^2 + 2x_1 x'_1 x_2 x'_2 \end{aligned}$$

Then $M = 6$, and if we choose


$$h_1(x) = 1, h_2(x) = \sqrt{2}x_1, h_3(x) = \sqrt{2}x_2, h_4(x) = x_1^2, h_5(x) = x_2^2,$$

$$\text{and } h_6(x) = \sqrt{2}x_1 x_2,$$

$$\text{then } K(x, x') = \langle h(x), h(x') \rangle.$$

Dim $h(x)$ infinite

SVM - Radial Kernel in Feature Space



- Fraction of support points depends on overlap; here 45%.
- The smaller B , the smaller the overlap, and more wiggly the function.
- B controls [generalization error](#).


Curse of Dimensionality

Support Vector Machines can suffer in high dimensions.

Method	Test Error (SE)	
	No Noise Features	Six Noise Features
1 SV Classifier	0.450 (0.003)	0.472 (0.003)
2 SVM/poly 2	0.078 (0.003)	0.152 (0.004)
3 SVM/poly 5	0.180 (0.004)	0.370 (0.004)
4 SVM/poly 10	0.230 (0.003)	0.434 (0.002)
5 BRUTO	0.084 (0.003)	0.090 (0.003)
6 MARS	0.156 (0.004)	0.173 (0.005)
Bayes	0.029	0.029

The addition of 6 noise features to the 4-dimensional feature space causes the performance of the SVM to degrade. The true decision boundary is the surface of a sphere, hence a quadratic monomial (additive) function is sufficient.

SVM via Loss + Penalty



With $f(x) = h(x)^T \beta + \beta_0$ and $y_i \in \{-1, 1\}$, consider

$$\min_{\beta_0, \beta} \sum_{i=1}^N [1 - y_i f(x_i)]_+ + \lambda \|\beta\|^2$$

Solution identical to SVM solution, with $\lambda = \lambda(B)$.

In general
$$\min_{\beta_0, \beta} \sum_{i=1}^N L[y_i, f(x_i)] + \lambda \|\beta\|^2$$

Loss Functions

For $Y \in \{-1, 1\}$

Log-likelihood: $L[Y, f(X)] = \log(1 + e^{-Yf(X)})$

- (negative) binomial log-likelihood or **deviance**.
- estimates the **logit**

$$f(X) = \log \frac{\Pr(Y = 1|X)}{\Pr(Y = -1|X)}$$

SVM: $L[Y, f(X)] = (1 - Yf(X))_+$.

- Called “**hinge loss**”
- Estimates the **classifier** (threshold)

$$C(x) = \text{sign} \left(\Pr(Y = 1|X) - \frac{1}{2} \right)$$

SVM and Function Estimation

SVM with general kernel K minimizes:

$$\sum_{i=1}^N (1 - y_i f(x_i))_+ + \lambda \|f\|_{\mathcal{H}_K}^2$$

with $f = b + h$, $h \in \mathcal{H}_K$, $b \in \mathcal{R}$. \mathcal{H}_K is the reproducing kernel Hilbert space (RKHS) of functions generated by the kernel K . The norm $\|f\|_{\mathcal{H}_K}$ is generally interpreted as a roughness penalty.

More generally we can optimize

$$\sum_{i=1}^N L(y_i, f(x_i)) + \lambda \|f\|_{\mathcal{H}_K}^2$$

Quadratic Programming (path algorithm)

$$L_P : \sum_{i=1}^N \xi_i + \frac{\lambda}{2} \beta^T \beta + \sum_{i=1}^N \alpha_i (1 - y_i f(x_i) - \xi_i) - \sum_{i=1}^N \gamma_i \xi_i$$

$$\frac{\partial}{\partial \beta} : \quad \beta = \frac{1}{\lambda} \sum_{i=1}^N \alpha_i y_i x_i$$

$$\frac{\partial}{\partial \beta_0} : \quad \sum_{i=1}^N y_i \alpha_i = 0,$$

along with the KKT conditions

$$\alpha_i (1 - y_i f(x_i) - \xi_i) = 0$$

$$\gamma_i \xi_i = 0$$


$$1 - \alpha_i - \gamma_i = 0$$

Implications of the KKT conditions

Observations are in one of three states:

- $\mathcal{L} = \{i : y_i f(x_i) < 1, \alpha_i = 1\}$, \mathcal{L} for Left of the elbow
 - $\mathcal{E} = \{i : y_i f(x_i) = 1, 0 \leq \alpha_i \leq 1\}$, \mathcal{E} for Elbow
 - $\mathcal{R} = \{i : y_i f(x_i) > 1, \alpha_i = 0\}$, \mathcal{R} for Right of the elbow
- Start with λ large, and the margin very wide. All $\alpha_i = 1$ (if $N_+ = N_-$). As $\lambda \downarrow 0$, the margin gets narrower.
 - For the narrowing margin to pass through a point, its α has to change from 1 to 0 (or from 0 to 1). While this is happening, the point has to **linger** on the margin. Hence the point moves from \mathcal{L} to \mathcal{R} via \mathcal{E} .
 - The condition $\sum_i y_i \alpha_i = 0$ demands a certain balance on opposite margins.




Example



- $\lambda = 0.5$, and the width of the soft margin is $2/\|\beta\| = 2 \times 0.587$.
- Two hollow points $\{3, 5\}$ are misclassified, while the two solid points $\{10, 12\}$ are correctly classified, but on the wrong side of their margin $f(x) = +1$; each of these has $\xi_i > 0$.
- The three square shaped points $\{2, 6, 7\}$ are exactly on the margin.


The Path

- The α_i are piecewise-linear in λ (or $1/C$)




- The points in \mathcal{E} characterize these paths, since points must stay on the margin ($y_i f(x_i) = 1$) while their α_i lie in $(0, 1)$.
- Points can revisit the margin more than once.
- The coefficients β_0 and β are piecewise-linear in $C = 1/\lambda$. (LARS, Efron et. al., 2002): quadratic criterion, L_1 constraint.
- The margins can stay wedged while their α_i change, if they are “loaded to capacity”.
- For non-separable data, the loss $\sum_i \xi_i$ achieves a minimum value, with a positive margin.


Piecewise Linear α Paths




Step: 14 Error: 2 Elbow Size: 3 Margin: 4.38



Path Statistics



The Need for Regularization



- γ is a kernel parameter: $K(x, z) = \exp(-\gamma \|x - z\|^2)$.
- λ (or C) are regularization parameters, which have to be determined using some means like cross-validation.

SVMs for regression

- Linear regression model:

$$f(x) = x^T \beta + \beta_0, \quad (1)$$



- To estimate β , we consider minimization of

$$H(\beta, \beta_0) = \sum_{i=1}^N V(y_i - f(x_i)) + \frac{\lambda}{2} \|\beta\|^2, \quad (2)$$

where

$$V_{\epsilon}(t) = \begin{cases} 0 & \text{if } |t| < \epsilon, \\ |t| - \epsilon, & \text{otherwise.} \end{cases} \quad (3)$$

This is called “ ϵ -insensitive” error measure.



The left panel shows the ϵ -insensitive error function used by the support vector regression machine. The right panel shows the error function used in Huber's robust regression (green curve). Beyond $|c|$, the function changes from quadratic to linear.

Software

- In R, `e1071` package and `library(svmpath)` available from CRAN.
- Many other packages fit SVMs